

AlexNet Experiments Results

Base Model:

1. Activation: ReLU
2. Dropout(p=0.5)
3. Pooling: Overlapping
4. Optimizer: SGD (with momentum and weight decay)

Training:


- **90 epochs** (validation accuracy measured after each epoch)
- Test Accuracy reported for model with **best validation accuracy**.

Note: The following results are w/o local response normalization. I also performed all the experiments with LRN, but they took more than twice the training time and didn't give any improvement. The results for the models with LRN are also included with the submission.

| Experiment | Training Accuracy (%) | Training Time (90 epochs) | Test Accuracy (%) |
|--------------------------------|-----------------------|---------------------------|-------------------|
| Base Model | 73.76 | 167m 1s | 72.89 |
| Base + Tanh | 67.17 | 163m 10s | 66.78 |
| Base + No Dropout | 82.15 | 184m 19s | 69.88 |
| Base + Non-overlapping pooling | 73.07 | 163m 18s | 71.75 |
| Base + SGD | 54.81 | 291m 36s | 61.18 |
| Base + SGD(with momentum) | 74.58 | 175m 9s | 72.73 |
| Base + Adam (lr=0.001) | 34.42 | 166m 60s | 43.14 |

1. Base Model

Base Model

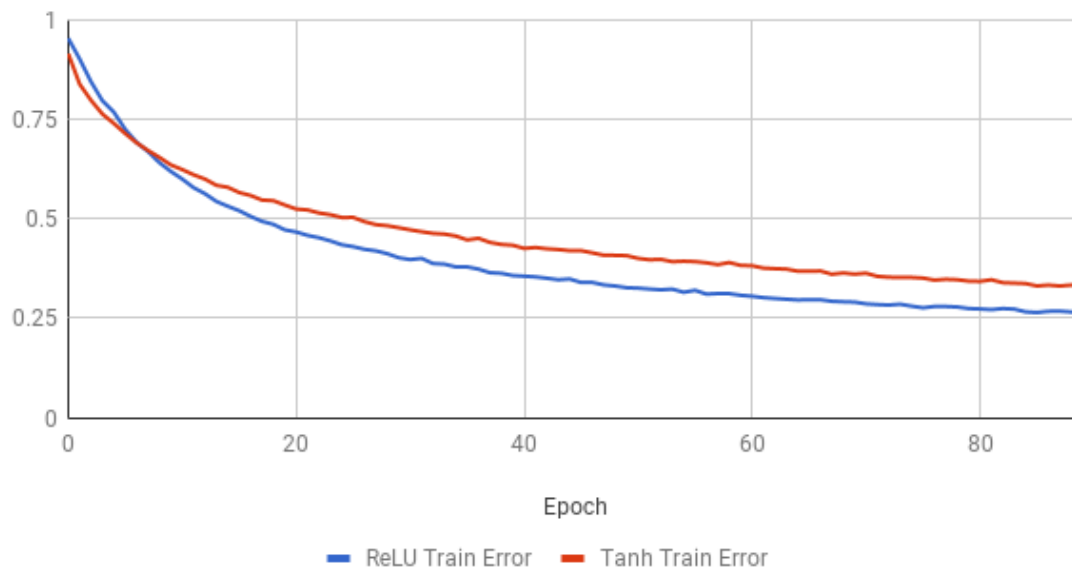


Test Time Per Example: ~2 ms

Final Test Error: 27.11%

2. Relu vs Tanh

ReLU vs Tanh



3. Dropout v/s No Dropout

| Type | Train Accuracy (%) | Best Validation Accuracy (%) | Test Accuracy (%) |
|-------------|--------------------|------------------------------|-------------------|
| w/ Dropout | 73.76 | 74.39 | 72.89 |
| w/o Dropout | 82.15 | 71.95 | 69.88 |


4. Overlapping Pooling v/s Non-Overlapping Pooling

| Type | Train Accuracy (%) | Best Validation Accuracy (%) | Test Accuracy (%) |
|-----------------|--------------------|------------------------------|-------------------|
| Overlapping | 73.76 | 74.39 | 72.89 |
| Non-overlapping | 73.07 | 73.41 | 71.75 |

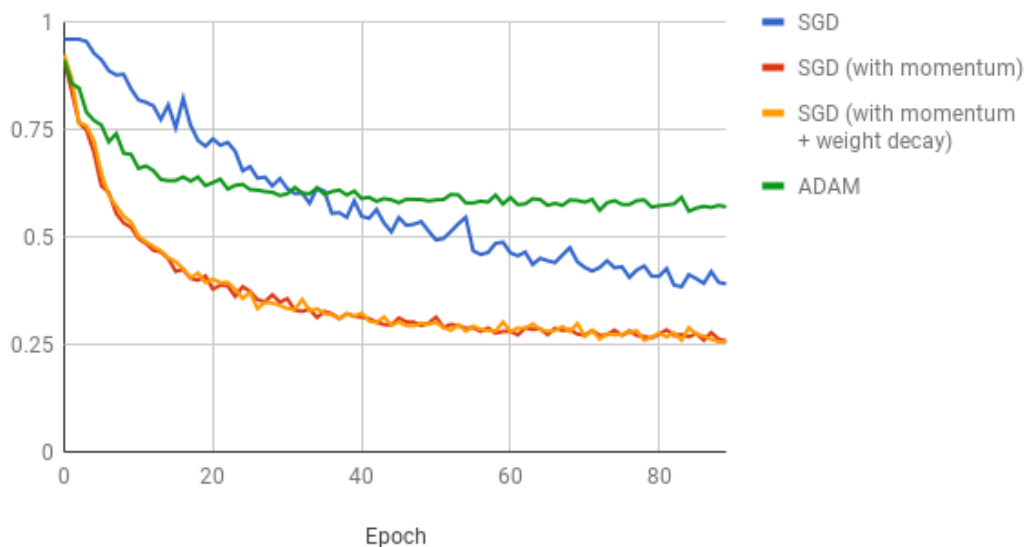
As can be observed, from the table, overlapping pooling clearly improves performance over non-overlapping pooling. According to the paper, it's because overlapping pooling makes it more difficult to overfit and hence improves test performance. I also found in an [reddit AMA](#) by Geoffrey Hinton, that non-overlapping pooling tends to lose positional information which is required to detect relationships between parts of an object and that overlapping allows some of this information to be preserved.

5. Optimization Techniques

Training Error with different Optimisation Techniques



Validation Error with different Optimisation techniques



6. Best Model

Validation Accuracy: 77.12%

Test Accuracy: 76.22%

Link: <https://drive.google.com/open?id=0By07sE0zY59RSE0RGF6STNzd1k>